
4.3:

Answer:

We can prove that the greedy algorithm currently in use actually minimizes the number of trucks that are needed using a similar analysis to that used for the *Interval Scheduling Problem*. We will establish the optimality of this greedy packing algorithm by identifying a measure under which it "stays ahead" of all other solutions. Namely, it "stays ahead" by packing more boxes into fewer trucks.

Suppose the greedy algorithm fits boxes b_1, b_2, \dots, b_j into the first k trucks. Now assume the other solution fits b_1, b_2, \dots, b_i into the first k trucks. We will induct on the number of trucks k to prove that $i \leq j$. Note this establishes the optimality of the greedy algorithm because it packs the most boxes into k trucks, thereby minimizing the number of trucks.

Basis Step:

$k = 1$. By definition, the greedy algorithm fits as many boxes as possible into the first truck. \Rightarrow basis step true.

Inductive Step:

Inductive Hypothesis: Assume the greedy algorithm packs j' boxes into the first $k - 1$ trucks, the other algorithm packs i' boxes into the first $k - 1$ trucks, and that $i' \leq j'$.

For the k th truck, the other solution packs in the additional boxes $b_{i'+1}, b_{i'+2}, \dots, b_i$ into the truck. But since $i' \leq j'$, we know that the greedy algorithm can fit at least $b_{j'+1}, b_{j'+2}, \dots, b_i$ additional boxes into the k th truck.

The result follows by the principle of mathematical induction and we have proved the optimality of the greedy algorithm currently in use by the trucking company.

4.7:

Answer:

We can prove using an exchange argument that the following polynomial-time algorithm finds a *schedule* with as small a *completion time* as possible:

Run the jobs J_1, J_2, \dots, J_n in order of decreasing finishing time f_i .

Let G be the resulting schedule. The exchange argument will prove that G is an optimal schedule. Note that regardless of the order, the last job is handed off to a high-end PC at the same time. Hence, this algorithm puts the job with the shortest finishing time as the last job in the resulting schedule G .

We want to prove that for any schedule $O \neq G$ that we can convert O into G without increasing the completion time, by swapping adjacent jobs to gradually modify O . Assume O has a different order than G . By the definition of G , O must contain an inversion. Namely, O must contain two jobs J_i and J_j such that J_i runs directly before J_j but $f_i < f_j$. We can optimize O by swapping J_i and J_j . Suppose O' is a schedule in which we swap J_i and J_j . Note that the finishing times of all jobs except J_i and J_j stay the same.

In O' , J_j now schedules earlier and will finish earlier than in O . The job J_i schedules later, but the supercomputer hands off J_i to a high-end PC in O' at the same time it would've handed off J_j in O . But the job J_i will finish earlier in O' than J_j would've finished in O . Hence O' does not have a greater completion time than O . By repeatedly performing such swaps on inversions and hence decreasing the number of inversions, we can therefore convert O into G without increasing the completion time. Hence the completion time for G is not greater than the completion time for any arbitrary schedule O , and G is an optimal schedule.

4.12:

Answer:**(a)**

This claim is false. We can prove this using a counterexample. Suppose we have stream 1 with $t_1 = 1$ and $b_1 = 50$, stream 2 with $t_2 = 1$ and $b_2 = 10$, and link parameter $r = 40$. Then the claim is false for stream 1. But there exists a valid schedule that runs the streams in the order stream 2, stream 1.

(b)

The algorithm to determine whether there exists a valid schedule can be described as follows. We assume that the streams are given and sent in increasing order of rate $r_i = \frac{b_i}{t_i}$. We claim that if the inequality $\sum_{i=1}^n b_i \leq r \sum_{i=1}^n t_i$ succeeds, the order gives a valid schedule. But if the inequality fails, no ordering produces a valid schedule. This is because in a total time of $\sum_{i=1}^n t_i$ we need to send $\sum_{i=1}^n b_i$ bits regardless of the order. We claim that the inequality also fails if the order sends too many bits for any initial time period $[0, t]$. Consider a time t and suppose i is the stream sent during the last time period. If $r_i \leq r$ then all streams sent so far have a rate at most r and the total sent is at most rt . But this contradicts the assumption that we sent too much in t time. So $r_i > r$. In any time period after t we also will send at least r bits, and therefore the total rate at the end of all streams also will break the rule of having an average rate of at most r . The algorithm runs in $O(n)$ time because we only need to check the one inequality $\sum_{i=1}^n b_i \leq r \sum_{i=1}^n t_i$ to determine if a valid schedule exists, which takes a linear amount of steps.

4.16:

Answer:

We can solve this problem using the following greedy algorithm. The strategy involves matching the n account events from the suspected bank account with the approximate time-stamps from the suspicious transactions. The approximate time-stamps can be expressed as a set of intervals $[t_i - e_i, t_i + e_i]$.

```

For i = 1, 2, ..., n
    If there are unmatched intervals containing  $x_i$ 
        Match  $x_i$  with the interval ending earliest
    Else
        Return there is no perfect matching
Endfor

```

If the greedy algorithm succeeds, then we have found a perfect matching. We can prove using an exchange argument that if there is a perfect matching, our algorithm will find it. To set up our proof by contradiction, assume there is a perfect matching but that our greedy algorithm does not find it. Suppose a perfect matching S contains matches for the account events x_1, x_2, \dots, x_i and i is the largest number with this property. Let's suppose the next account event x_{i+1} is matched to the interval around t_l , but our greedy algorithm says x_{i+1} should match to the interval around t_j . By definition of the greedy algorithm, $t_j + e_j \leq t_l + e_l$. But assume t_j is matched to x_k in S , where $x_k \geq x_{i+1}$. Thus we know $t_l - e_l \leq x_{i+1} \leq x_k \leq t_j + e_j \leq t_l + e_l$. But this means we can do a swap and instead match x_k with the interval around t_l and x_{i+1} with the interval around t_j , resulting in a new perfect matching S' that agrees with our greedy algorithm. But S' agrees with our greedy algorithm on the first $i + 1$ account events. $\Rightarrow \Leftarrow$ The running time of this algorithm is $O(n^2)$, as required, because there are n iterations of the for loop and the step "Match x_i with the interval ending earliest" takes $O(n)$ time to iterate over the unmatched intervals.